

Rishabh Gupta

Stony Brook, NY | +91 9773744566 | rishabh.gupta.1@stonybrook.edu | [GitHub](#) | [LinkedIn](#)

EDUCATION

Stony Brook University

Master of Science in Computer Science

Stony Brook, New York

Expected Graduation: Dec. 2027

Relevant Coursework: System Security, Data Science, Advancements in Artificial Intelligence, Natural Language

LNM Institute of Information Technology

Bachelor of Science in Computer Science

Jaipur, India

Relevant Coursework: Computer Networks, Data Structures, Python Programming, Object Oriented Programming

Aug. 2018 – July 2022

TECHNICAL SKILLS

Primary: 4 years- Python, FastAPI, SQL, Cassandra, RAG, DSPy, LangChain, MCP, Git, REST APIs, React, Redux

Additional: 4 years- Kubernetes, Docker, Aws, Jenkins, LLD, Data Structures and Algorithms, Springboot, JavaScript

Practices: 4 years- Object Oriented Programming (OOPs), Agile, Effective communication

PROFESSIONAL EXPERIENCE

Rocket Software

Hyderabad, India

Software Engineer II- AI

Sept. 2024 – Present

- Developed metadata-aware query router for [SmartChat](#) a RAG based System using DSPy and automated SQL generation, improving answer quality by 84% over baseline semantic/hybrid retrieval on complex customer data.
- Implemented unified LLM and embedding inference layer abstracting multiple provider APIs (OpenAI, Anthropic, AWS Bedrock), reducing time to onboard new models and cutting documentation overhead.
- Built a centralized runtime config system with SSE-based pub-sub and Gunicorn SIGHUP reloads, enabling live config updates across services without Kubernetes pod restarts.
- Executed 200+ unit, integration, and functional test cases, elevating code coverage from 70% to 95% and improving system reliability.

Capgemini India

Pune, India

Software Developer

Oct. 2022 – Sept. 2024

- Architected trie + priority queue for Data Catalog routing, slashed query latency 65% across 200k+ records.
- Devised C++ indexing microservice (gRPC to FastAPI), accelerated 15k docs/sec at 40ms p99 for hybrid search
- Delivered WCAG 2.1 catalog dashboard (ARIA, keyboard nav, NVDA tested) for accessible metadata browsing.

Capgemini India

Pune, India

Software Developer Intern – Data Governance

March 2022 – May 2022

- Optimized Alation lineage service with tracemalloc profiling, cutting memory 35% via bounded caches and weakref cleanup for 100k+ assets.
- Scaled service to slash governance audit times 40%, accelerating compliance reporting across enterprise data.

PUBLICATIONS & RECOGNITIONS

Hackathons: Smart Scrum Master- [Most Innovative project Award](#) at RocketBuild 2025.

Boston, MA

- Automated Jira ticket creation from Microsoft Teams meetings using a custom MCP server and AI agent, achieving 90%+ accuracy and reducing manual backlog grooming effort by 70%.
- Boosted sprint planning efficiency by 40% through a Jira MCP AI agent that enabled seamless ticket creation, updates, and prioritization directly from meeting insights.
- Reduced sprint review preparation time by 50% by leveraging AI-driven analytics to deliver real-time insights on team velocity, burndown and blockers.